

A S T R E A

AI-Powered Personal Reminder Assistant

PROPOSAL

The Problem

Managing reminders should be effortless, yet today's tools force users into rigid, manual workflows that break the moment life gets complex.

Creating a simple reminder like "call mom tomorrow at 3pm" should take seconds, but most apps make you tap through date pickers, time selectors, and category menus just to capture a thought. Users who switch between Android and iOS have it worse: they lose their entire reminder system every time they change phones because nothing truly syncs across ecosystems. When a notification finally fires, the only options are "dismiss" or "open the app." There's no way to snooze for a custom duration, mark something complete, or take any meaningful action without leaving what you're doing. On top of that, recurring rules barely go beyond "daily" or "weekly," and asking something like "what do I have related to health this month?" returns nothing because these apps match keywords, not meaning.

The core issue is that people think in natural language. They say "remind me to pick up the prescription after work" and expect their app to just get it. Instead, they're forced to adapt to the software. Astrea flips that: the software adapts to the user.

Target Audience

Astrea is built for people like Sam Beckman, someone who lives by reminders and uses them for everything from daily errands to complex recurring schedules. Sam switches between Android and iOS and is frustrated that no app truly syncs across platforms. He wants reminders that sync on actions: when he dismisses, snoozes, or completes on one device, it reflects everywhere instantly. He needs custom snooze durations straight from notifications, powerful recurring rules, and an assistant that understands what he means, not just what he types.

More broadly, Astrea serves platform switchers who need true cross-device sync, productivity enthusiasts who need powerful recurring rules and smart categorization, voice-first users who prefer speaking over typing, and overwhelmed organizers who need an assistant that retrieves reminders by meaning. The reminder market is valued at over \$4.5B globally, yet tools like Apple Reminders, Google Tasks, and Todoist still rely on form-based input. None offer natural language understanding as the primary interface. Astrea fills that gap.

Monetization Strategy

Astrea follows a freemium model where core reminder functionality is free and unlimited, while AI-powered features are gated behind a subscription. Users experience real value before being asked to pay, and the upgrade path feels natural because they hit the limit precisely when they're most engaged.

Plans

	Free	Astrea Pro
AI messages	25 per day	400 per day
Reminder management	Unlimited	Unlimited
Cross-device sync	Included	Included
Push notifications	Included	Included
Custom snooze from notifications	Included	Included
Recurring reminder rules	Included	Included
Semantic search (ask questions)	Limited	Unlimited
Price	\$0	\$4.99/mo or \$39.99/yr

This model works because the free plan is fully functional, so users never feel locked out of core features. The daily AI limit is reached when users are most active, creating a high-intent upgrade moment. By that point they've already experienced AI-powered reminders firsthand, so the value is proven before they're asked to pay. Monthly and yearly subscription options are managed through RevenueCat with a built-in Customer Center for self-service management, which helps reduce churn.

Revenue Infrastructure

Component	Technology
Payment processing	Google Play Billing / Apple IAP
Subscription management	RevenueCat SDK
Paywall UI	RevenueCat remote paywalls (A/B testable from dashboard)
Entitlement	Astrea Pro single entitlement, checked client-side and server-side
Analytics	RevenueCat dashboard for MRR, churn, and trial conversion

With a **2 to 5% free-to-paid conversion rate**, which is industry standard for productivity apps, and an average revenue per paying user of roughly **\$45 per year**, Astrea reaches sustainable revenue at modest scale. The AI cost per user stays low because LLM calls are per-message rather than per-user, meaning margins improve as the user base grows.

